

LabAssignment2_Optimization

June 11, 2023

1 Lab Assignment 2

1.1 Breast Cancer Prediction

1. Predict Breast Cancer using Feed forward neural network of configuration Nx128x64x32x16x1 and evaluated performance of momemntum, Adagrad, RMSprop, and Adam optimizers in terms of accuracy and F1-score.

```
[2]: import pandas as pd
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, f1_score
import tensorflow as tf
from tensorflow.keras import models, layers
```

```
[3]: df = pd.read_csv("data/BreastCancerData.csv")
```

```
[4]: df.head()
```

```
[4]:
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	\
0	842302	M	17.99	10.38	122.80	1001.0	
1	842517	M	20.57	17.77	132.90	1326.0	
2	84300903	M	19.69	21.25	130.00	1203.0	
3	84348301	M	11.42	20.38	77.58	386.1	
4	84358402	M	20.29	14.34	135.10	1297.0	

	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	\
0	0.11840	0.27760	0.3001	0.14710	
1	0.08474	0.07864	0.0869	0.07017	
2	0.10960	0.15990	0.1974	0.12790	
3	0.14250	0.28390	0.2414	0.10520	
4	0.10030	0.13280	0.1980	0.10430	

...	texture_worst	perimeter_worst	area_worst	smoothness_worst	\
0	17.33	184.60	2019.0	0.1622	
1	23.41	158.80	1956.0	0.1238	
2	25.53	152.50	1709.0	0.1444	
3	26.50	98.87	567.7	0.2098	

4	...	16.67	152.20	1575.0	0.1374
---	-----	-------	--------	--------	--------

	compactness_worst	concavity_worst	concave points_worst	symmetry_worst	\
0	0.6656	0.7119	0.2654	0.4601	
1	0.1866	0.2416	0.1860	0.2750	
2	0.4245	0.4504	0.2430	0.3613	
3	0.8663	0.6869	0.2575	0.6638	
4	0.2050	0.4000	0.1625	0.2364	

	fractal_dimension_worst	Unnamed: 32
0	0.11890	NaN
1	0.08902	NaN
2	0.08758	NaN
3	0.17300	NaN
4	0.07678	NaN

[5 rows x 33 columns]

```
[5]: df.shape
```

```
[5]: (569, 33)
```

```
[6]: df.isna().sum()
```

```
[6]: id                0
      diagnosis         0
      radius_mean       0
      texture_mean       0
      perimeter_mean     0
      area_mean          0
      smoothness_mean    0
      compactness_mean   0
      concavity_mean     0
      concave points_mean 0
      symmetry_mean      0
      fractal_dimension_mean 0
      radius_se          0
      texture_se         0
      perimeter_se       0
      area_se            0
      smoothness_se      0
      compactness_se     0
      concavity_se       0
      concave points_se  0
      symmetry_se        0
      fractal_dimension_se 0
      radius_worst       0
```

```

texture_worst            0
perimeter_worst          0
area_worst               0
smoothness_worst         0
compactness_worst        0
concavity_worst          0
concave points_worst     0
symmetry_worst           0
fractal_dimension_worst  0
Unnamed: 32              569
dtype: int64

```

```
[7]: df=df.drop(columns=['Unnamed: 32','id'])
```

```
[8]: df.head()
```

```

[8]:  diagnosis  radius_mean  texture_mean  perimeter_mean  area_mean  \
0          M         17.99         10.38         122.80      1001.0
1          M         20.57         17.77         132.90      1326.0
2          M         19.69         21.25         130.00      1203.0
3          M         11.42         20.38          77.58       386.1
4          M         20.29         14.34         135.10      1297.0

      smoothness_mean  compactness_mean  concavity_mean  concave points_mean  \
0          0.11840         0.27760         0.3001         0.14710
1          0.08474         0.07864         0.0869         0.07017
2          0.10960         0.15990         0.1974         0.12790
3          0.14250         0.28390         0.2414         0.10520
4          0.10030         0.13280         0.1980         0.10430

      symmetry_mean  ...  radius_worst  texture_worst  perimeter_worst  \
0          0.2419  ...         25.38         17.33         184.60
1          0.1812  ...         24.99         23.41         158.80
2          0.2069  ...         23.57         25.53         152.50
3          0.2597  ...         14.91         26.50          98.87
4          0.1809  ...         22.54         16.67         152.20

      area_worst  smoothness_worst  compactness_worst  concavity_worst  \
0         2019.0         0.1622         0.6656         0.7119
1         1956.0         0.1238         0.1866         0.2416
2         1709.0         0.1444         0.4245         0.4504
3          567.7         0.2098         0.8663         0.6869
4         1575.0         0.1374         0.2050         0.4000

      concave points_worst  symmetry_worst  fractal_dimension_worst
0          0.2654         0.4601         0.11890
1          0.1860         0.2750         0.08902


```

2	0.2430	0.3613	0.08758
3	0.2575	0.6638	0.17300
4	0.1625	0.2364	0.07678

[5 rows x 31 columns]

```
[9]: df['diagnosis'].value_counts().plot(kind='bar')
```

```
[9]: <Axes: xlabel='diagnosis'>
```



```
[10]: encoder = LabelEncoder()
```

```
[11]: encoded_diagnosis = encoder.fit_transform(df['diagnosis'])
```

```
[12]: df['encoded_diagnosis']=encoded_diagnosis
```

```
[13]: df.head()
```

```
[13]:
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	\
0	M	17.99	10.38	122.80	1001.0	
1	M	20.57	17.77	132.90	1326.0	

2	M	19.69	21.25	130.00	1203.0
3	M	11.42	20.38	77.58	386.1
4	M	20.29	14.34	135.10	1297.0

	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	\
0	0.11840	0.27760	0.3001	0.14710	
1	0.08474	0.07864	0.0869	0.07017	
2	0.10960	0.15990	0.1974	0.12790	
3	0.14250	0.28390	0.2414	0.10520	
4	0.10030	0.13280	0.1980	0.10430	

	symmetry_mean	...	texture_worst	perimeter_worst	area_worst	\
0	0.2419	...	17.33	184.60	2019.0	
1	0.1812	...	23.41	158.80	1956.0	
2	0.2069	...	25.53	152.50	1709.0	
3	0.2597	...	26.50	98.87	567.7	
4	0.1809	...	16.67	152.20	1575.0	

	smoothness_worst	compactness_worst	concavity_worst	concave points_worst	\
0	0.1622	0.6656	0.7119	0.2654	
1	0.1238	0.1866	0.2416	0.1860	
2	0.1444	0.4245	0.4504	0.2430	
3	0.2098	0.8663	0.6869	0.2575	
4	0.1374	0.2050	0.4000	0.1625	

	symmetry_worst	fractal_dimension_worst	encoded_diagnosis
0	0.4601	0.11890	1
1	0.2750	0.08902	1
2	0.3613	0.08758	1
3	0.6638	0.17300	1
4	0.2364	0.07678	1

[5 rows x 32 columns]

```
[14]: df=df.drop(columns=['diagnosis'])
```

```
[15]: X=df.drop(columns=['encoded_diagnosis'])
y=df[['encoded_diagnosis']]
```

```
[16]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
```

```
[17]: def create_model():
    model = models.Sequential()
    model.add(layers.Dense(128, activation='relu', input_shape=(30,)))
    model.add(layers.Dense(64, activation='relu'))
```

```

model.add(layers.Dense(32, activation='relu'))
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
return model

```

```

[18]: optimizers = ['sgd', 'adagrad', 'rmsprop', 'adam']
      results = []

      for optimizer in optimizers:
          model = create_model()
          model.compile(optimizer=optimizer, loss='binary_crossentropy',
↪ metrics=['accuracy'])
          model.fit(X_train, y_train, epochs=10, batch_size=32, verbose=0)

          # Predict on the test set
          y_pred = model.predict(X_test)
          y_pred_class = (y_pred > 0.5).astype(int)
          # Calculate accuracy and F1-score
          accuracy = accuracy_score(y_test, y_pred_class)
          f1 = f1_score(y_test, y_pred_class)

          results.append({'Optimizer': optimizer, 'Accuracy': accuracy, 'F1-score':
↪ f1})

```

```

4/4 [=====] - 0s 6ms/step
4/4 [=====] - 0s 5ms/step
4/4 [=====] - 0s 5ms/step
4/4 [=====] - 0s 5ms/step

```

```

[19]: results_df = pd.DataFrame(results)
      results_df

```

```

[19]:   Optimizer  Accuracy  F1-score
0      sgd    0.622807  0.000000
1  adagrad    0.929825  0.900000
2  rmsprop    0.947368  0.926829
3      adam    0.912281  0.893617

```

From the result above we can see that the `rmsprop` optimizer performed best on the basis of F1-score and Accuracy

1.2 News Category Classification

2. Predict class label to predict news category using Feed forward neural network of configuration `Nx256x128x64x32x16x1` and evaluate performance of momentum, Adagrad, RMSprop, and Adam optimizers in terms of accuracy and F1-score.

There are four distinct sections where each story may fall in to. The Sections are labelled as follows

: * Politics: 0 * Technology: 1 * Entertainment: 2 * Business: 3

```
[20]: import pandas as pd
      from sklearn.preprocessing import LabelEncoder
      from sklearn.model_selection import train_test_split
      from sklearn.metrics import accuracy_score, f1_score
      import tensorflow as tf
      from tensorflow.keras import models, layers
      from tensorflow.keras.preprocessing.text import Tokenizer
      from tensorflow.keras.preprocessing.sequence import pad_sequences
      import numpy as np
```

```
[21]: df=pd.read_csv("data/NewsData.csv",encoding='latin')
```


```
[22]: df.head()
```

```
[22]:
```

	STORY	SECTION
0	But the most painful was the huge reversal in ...	3
1	How formidable is the opposition alliance amon...	0
2	Most Asian currencies were trading lower today...	3
3	If you want to answer any question, click on ...	1
4	In global markets, gold prices edged up today ...	3

```
[23]: df['SECTION'].value_counts().plot(kind='bar')
```

```
[23]: <Axes: xlabel='SECTION'>
```



```
[24]: X=df['STORY'].values
      y=df['SECTION'].values
```

```
[25]: # Split the data into training and testing sets
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪random_state=42)
```

```
[26]: labels = tf.keras.utils.to_categorical(y_train, num_classes=4)
```

```
[27]: # Tokenize the text data
      tokenizer = Tokenizer()
      tokenizer.fit_on_texts(X_train)
      word_index = tokenizer.word_index
      sequences = tokenizer.texts_to_sequences(X_train)
```

```
[28]: # Pad sequences to a fixed length
      max_sequence_length = 10 # maximum length of a sequence
      padded_sequences = pad_sequences(sequences, maxlen=max_sequence_length)
```

```
[29]: def create_model():
      model = tf.keras.models.Sequential([
```



```

    tf.keras.layers.Embedding(len(word_index) + 1, 256,
    ↪input_length=max_sequence_length),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(32, activation='relu'),
    tf.keras.layers.Dense(16, activation='relu'),
    tf.keras.layers.Dense(4, activation='softmax')
])
return model

```

```

[31]: optimizers = ['sgd', 'adagrad', 'rmsprop', 'adam']
      results = []

      for optimizer in optimizers:
          model = create_model()
          model.compile(optimizer=optimizer, loss='categorical_crossentropy',
          ↪metrics=['accuracy'])
          model.fit(padded_sequences, labels, epochs=10, batch_size=32, verbose=0)

          test_sequences = tokenizer.texts_to_sequences(X_test)
          test_padded_sequences = pad_sequences(test_sequences,
          ↪maxlen=max_sequence_length)
          # Predict on the test set
          y_pred = model.predict(test_padded_sequences)
          y_pred_class=[]
          for item in y_pred:
              temp = np.argmax(item)
              y_pred_class.append(temp)
          # Calculate accuracy and F1-score
          accuracy = accuracy_score(y_test, y_pred_class)
          f1 = f1_score(y_test, y_pred_class,average='macro')

          results.append({'Optimizer': optimizer, 'Accuracy': accuracy, 'F1-score':
          ↪f1})

```

```

48/48 [=====] - 1s 6ms/step
48/48 [=====] - 1s 7ms/step
48/48 [=====] - 1s 9ms/step
48/48 [=====] - 1s 14ms/step

```

```

[32]: results_df = pd.DataFrame(results)
      results_df

```

```

[32]:   Optimizer  Accuracy  F1-score
0      sgd      0.359764  0.135899
1  adagrad      0.359764  0.132289

```

2	rmsprop	0.699869	0.681804
3	adam	0.748362	0.735969

For News Classification adam performed better.